


Processes

Introduction: Von Neuman Model

- Both program and data reside in memory
- Execution stages in CPU:
 - Fetch instruction
 - Decode instruction
 - Execute instruction
 - Write back result
- OS is just a program
 - OS code and data reside in memory too
 - Invoke OS functionality through system calls



Execution Mode


- Two modes of execution for protection reasons
 - Privileged - kernel-mode
 - Non-privileged - user-mode
- Kernel executes in kernel-mode
 - Access hardware resources
 - Protected from interference by user programs
 - Portion of the OS
- User code executes in user-mode
- OS functionality that does not need direct access to hardware may run in user-mode
 - Microkernel design basis

Interrupts and traps

- **Interrupt: an asynchronous event**
 - External event
 - Independent instruction execution in the processor
 - E.g. DMA completion
 - Can be masked (specifically or not)
- **Trap: a synchronous software event**
 - Synchronous event
 - Caused by the execution of the current instruction
 - E.g. system calls, floating point error
 - Conditional or unconditional

More Interrupts and Traps

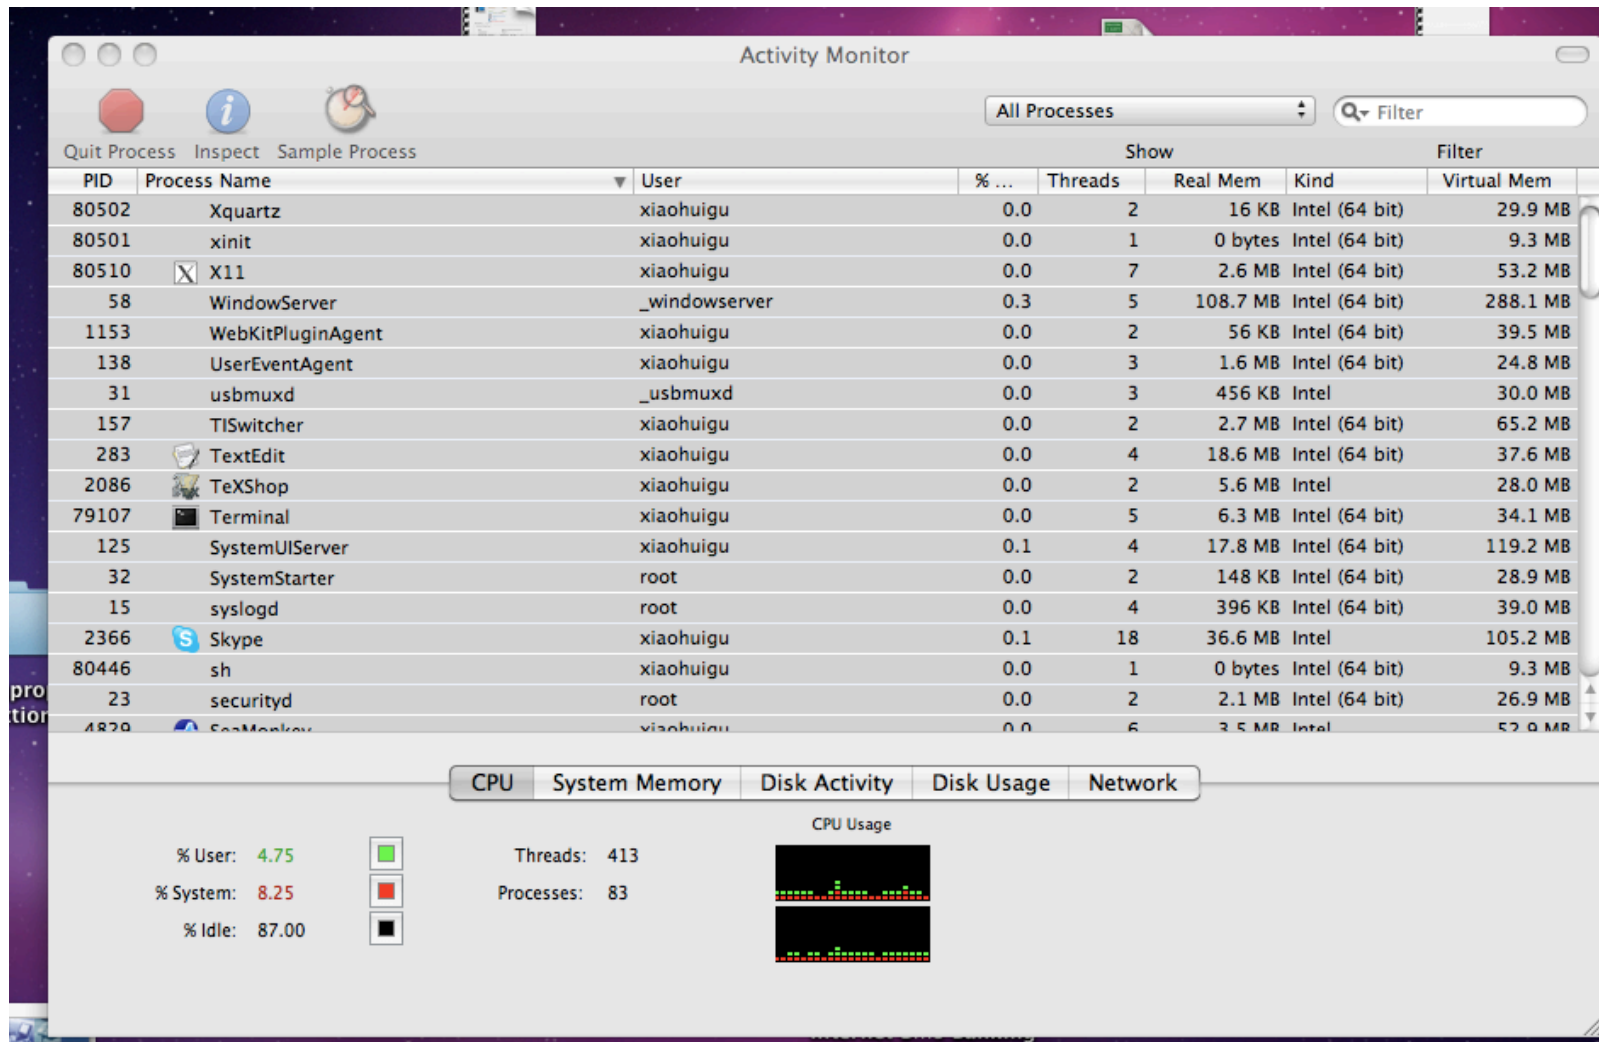
- **Interrupt and trap events**
 - Statically defined (typically as integers)
 - Each interrupt and trap has an associated interrupt vector
 - Interrupt vector specifies handler
 - Code that should be called when the event occurs
- **At interrupt or trap, processor**
 - Saves current state of execution
 - Jumps to the handler




Process

- Process – a program in execution; an “instantiation” of a program
- A process includes:
 - program counter
 - stack
 - data section

Processes



Process in Memory



Process Control Block (PCB)

Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information

Example PCB in XINU

/* excerpt from file proc.h */

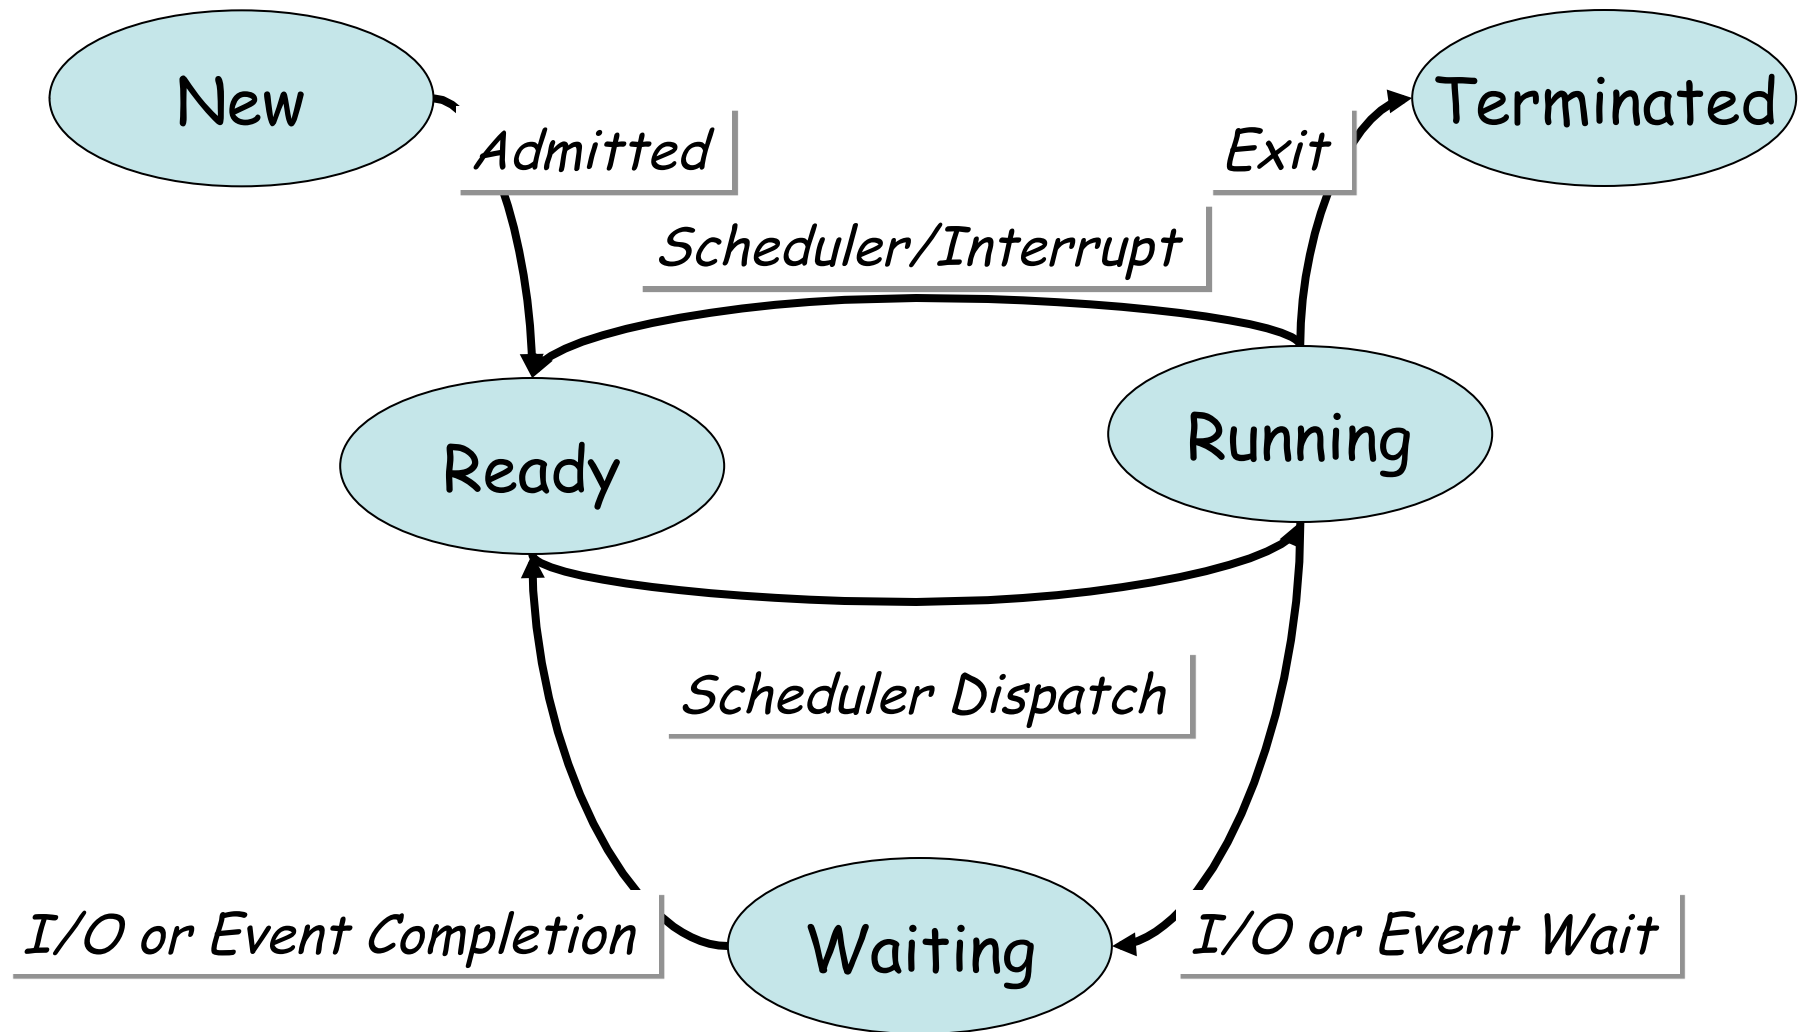
```
struct pentry {                                /* process table entry */
    char    pstate;                            /* process state: PRCURR, etc. */
    int     pprio;                             /* process priority */
    int     pesp;                              /* saved stack pointer */
    STATWORD pirmask;                          /* saved interrupt mask */
    int     psem;                              /* semaphore if process waiting */
    WORD    pmsg;                              /* message sent to this process */
    char    phasmg;                            /* nonzero iff pmsg is valid */
    WORD    pbase;                             /* base of run time stack */
    int     pstklen;                           /* stack length */
    WORD    plimit;                            /* lowest extent of stack */
    char    pname[PNMLEN];                    /* process name */
    int     pargs;                             /* initial number of arguments */
    WORD    paddr;                             /* initial code address */
    short   pdevs[2];                          /* devices to close upon exit */
    int     fildes[_NFILE];                   /* file - device translation */
};
```

```
extern struct pentry proctab[];
extern int     numproc;                       /* currently active processes */
extern int     nextproc;                     /* search point for free slot */
extern int     curripid;                     /* currently executing process */
```

Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution

Process Lifecycle

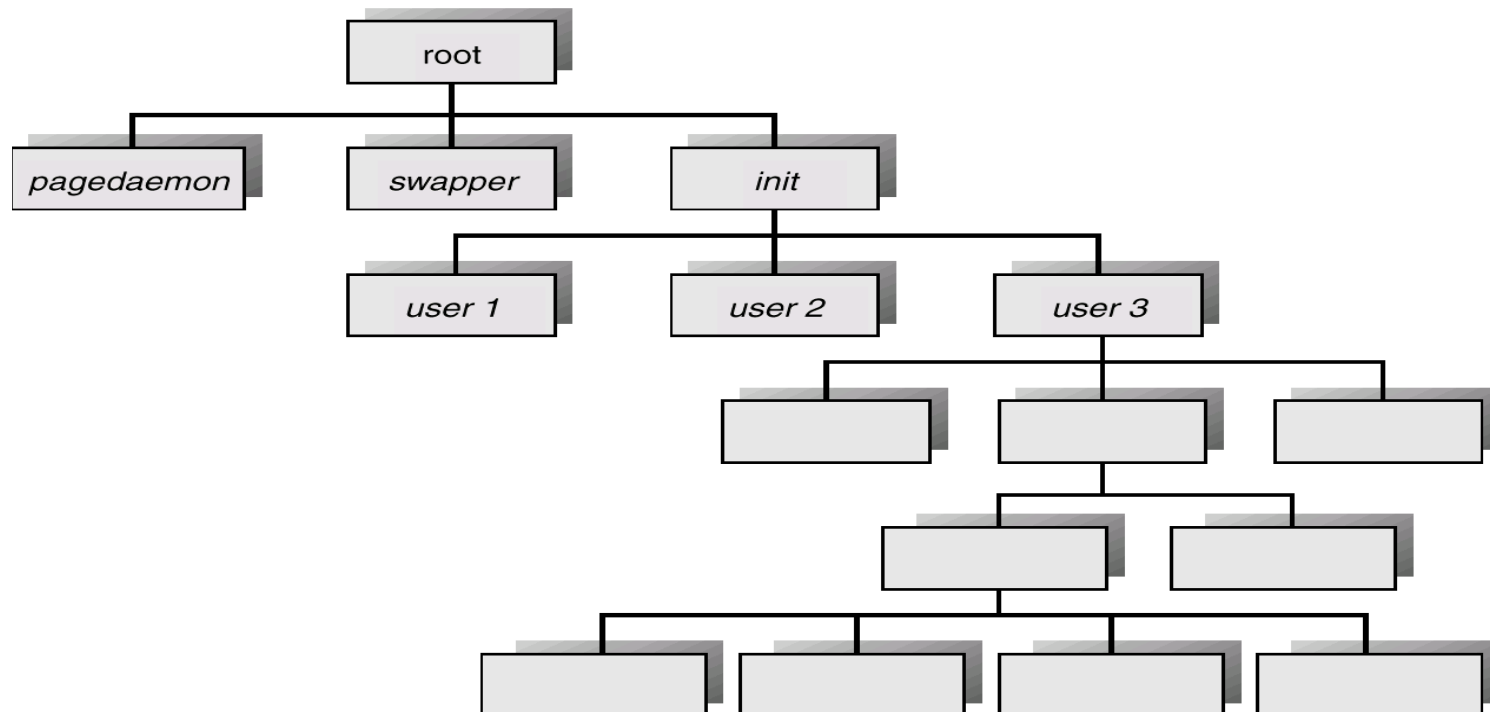


Process Manipulation

- Performed by OS routines
- Example operations
 - Creation
 - Termination
 - Suspension
 - Resumption
- State variable in process table records activity

Process Creation

- Parent process creates children processes,
 - Which, in turn create other processes,
 - Forming a tree of processes



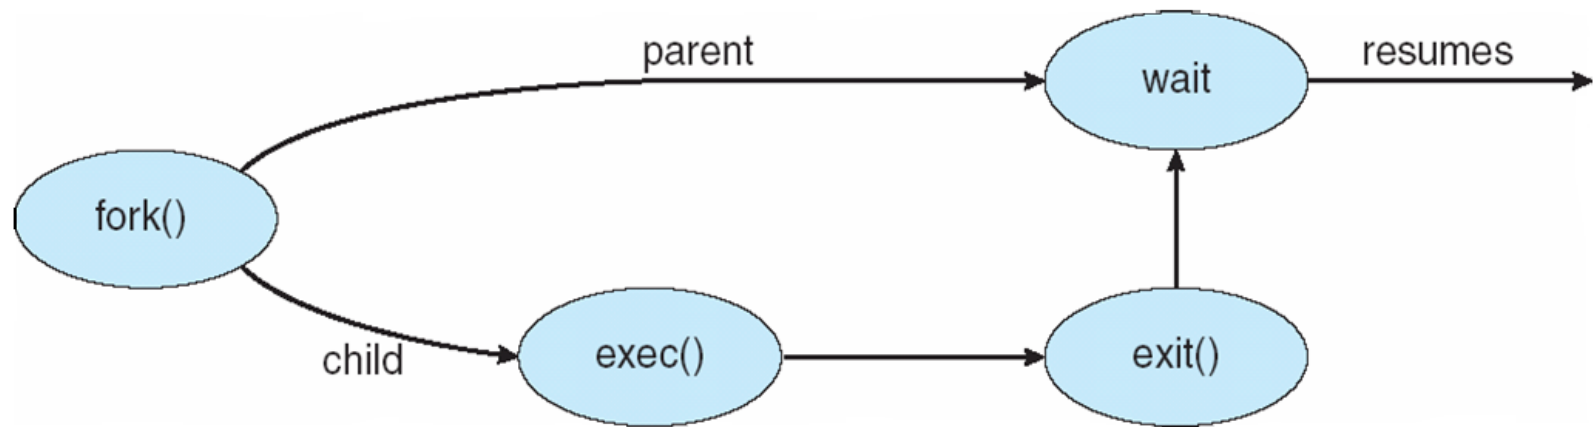
Process Creation

- Policy on resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Policy on execution
 - Parent and children execute concurrently
 - Parent waits until children terminate
- Policy on address space
 - Child duplicate of parent
 - Child has a program loaded into it

Process Creation (Cont.)

- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace new process' memory space with a new program

Process Creation



C Program Forking Separate Process

```
int main()
{
    pid_t  pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to
        complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

Process Termination

- Possible scenarios for process termination
 - Exit (by itself)
 - Abort (by parent)
 - Kill (by sysadmin)
- Exit
 - Process executes last statement and asks operating system to delete

Process Termination

- Abort
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - Some operating system do not allow child to continue if its parent terminates
 - All children terminated - ***cascading termination***
- Kill
 - Administration purpose

Process Suspension

- Temporarily “stop” a process
 - Prohibit from using the CPU
- Why?
- What should be done?
 - Change its state in PCB
 - Save its machine states for later resumption
 - Process table entry retained
 - Complete state saved

Context Switch

- When CPU switches to another process
- System must
 - Save the state of the old process (suspend) and
 - Load the saved state for the new process (resume)
- Context-switch time is overhead
 - System does no useful work while switching
- Time dependent on hardware support

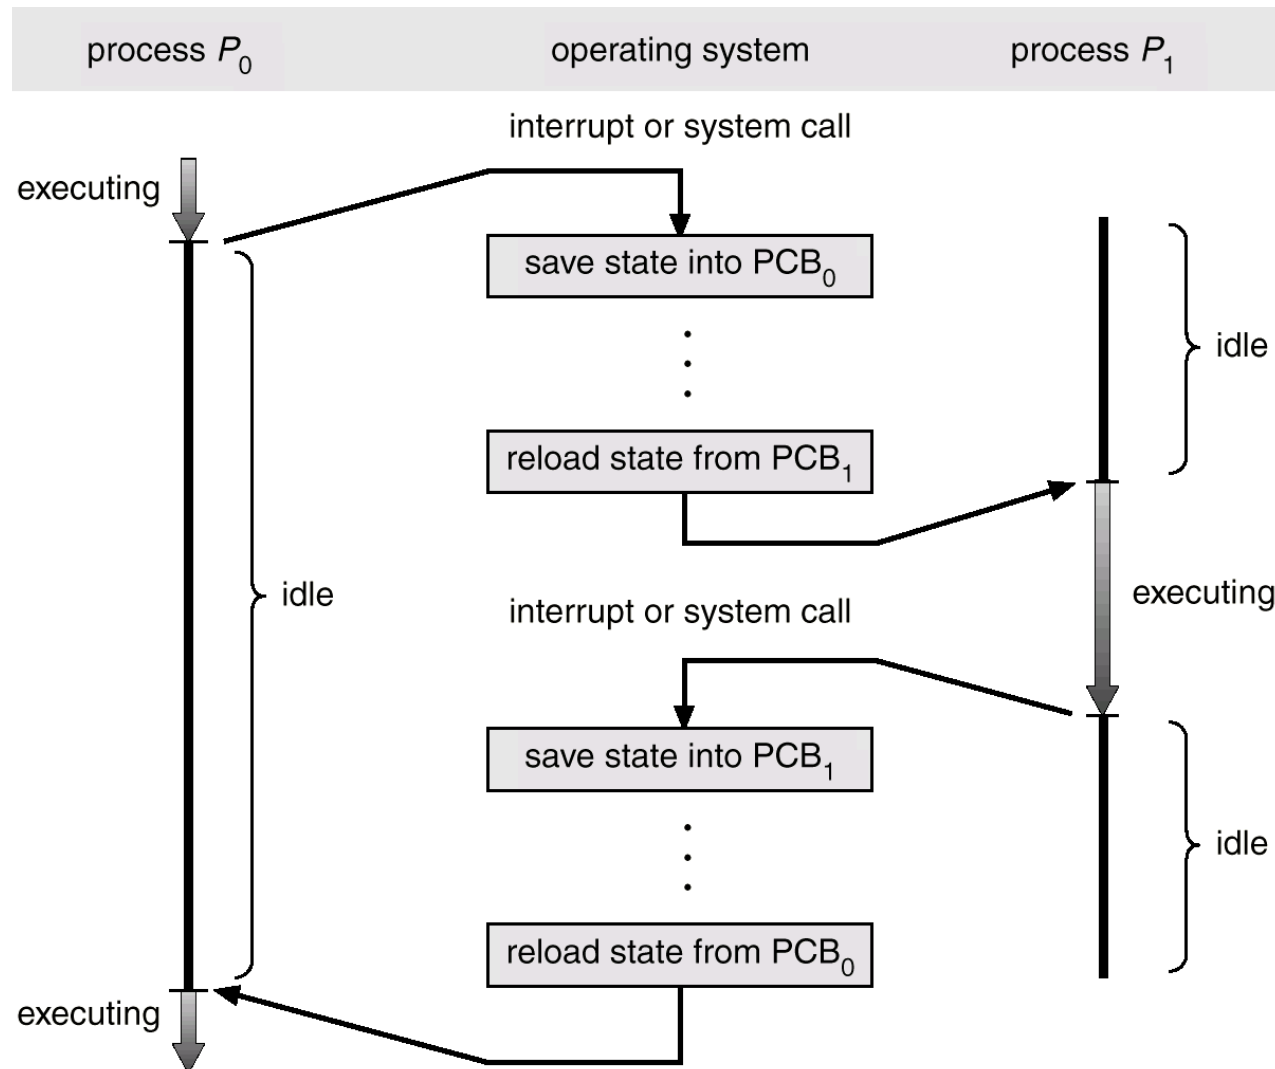
Context Switching

- How to do a context switch?
 - Very carefully!!
- Save state of currently executing process
 - Copy all “live” registers to process control block
 - Need at least 1 scratch register -- points to area of memory in process control block that registers should be saved to
- Restore state of process to run next
 - Copy values of live registers from process control block to registers
- How to get into and out of the context switching code?

Context Switching

- OS is just code stored in memory, so ...
 - Call context switching subroutine
 - The subroutine saves context of current process, restores context of the next process to be executed, and returns
 - The subroutine returns in the context of another (the next) process!
 - Eventually, will switch back to the current process
 - To process, it appears as if the context switching subroutine just took a long while to return

CPU Switch From Process to Process



Xinu Implementation

- Read relevant source code in Xinu
 - Process queue management
 - h/q.h sys/queue.c sys/insert.c, ...
 - Proc. creation/suspension/resumption/termination:
 - sys/create.c, sys/suspend.c sys/resume.c, sys/kill.c
 - Process scheduling
 - sys/resched.c
 - Other initialization code

Next Lecture

- Thread