



Threads

Chapter 4

Single and Multithreaded Processes

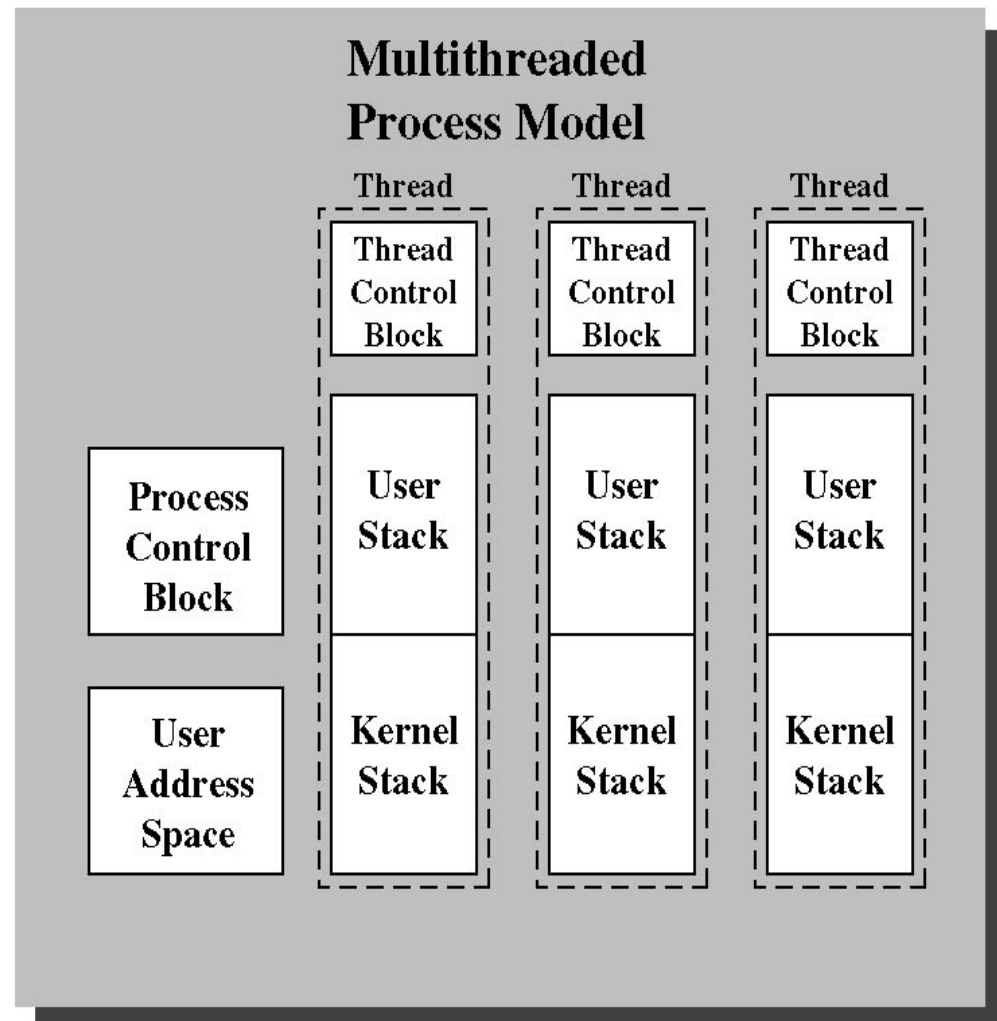
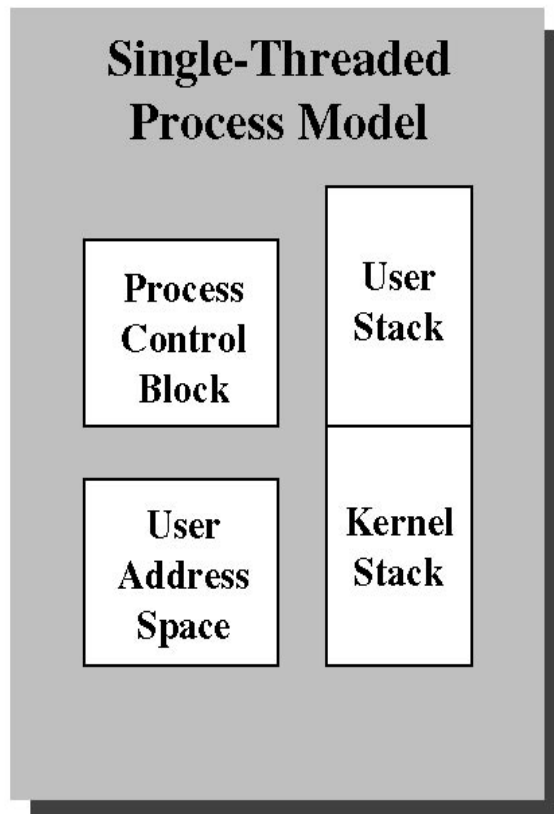


single-threaded process




multithreaded process


Single and Multithreaded Processes



Process Address Space



single threaded address space



multi-threaded address space

Thread

- In single threaded systems, a **process** is:
 - Resource owner
 - memory address space, files, I/O resources
 - Scheduling/execution unit
 - execution state/context, dispatch unit
- Multithreaded systems
 - Separation of resource ownership & execution unit
 - A **thread** is unit of execution, scheduling and dispatching
 - A **process** is a container of resources, and a collection of threads

Thread

- All threads of a process share resources
 - Memory address space: global data, code, heap ...
 - Open files, network sockets, other I/O resources
 - User-id
 - IPC facilities
- Private state of each thread:
 - Execution state: running, ready, blocked, etc..
 - Execution context: Program Counter, Stack Pointer, other user-level registers
 - Per-thread stack

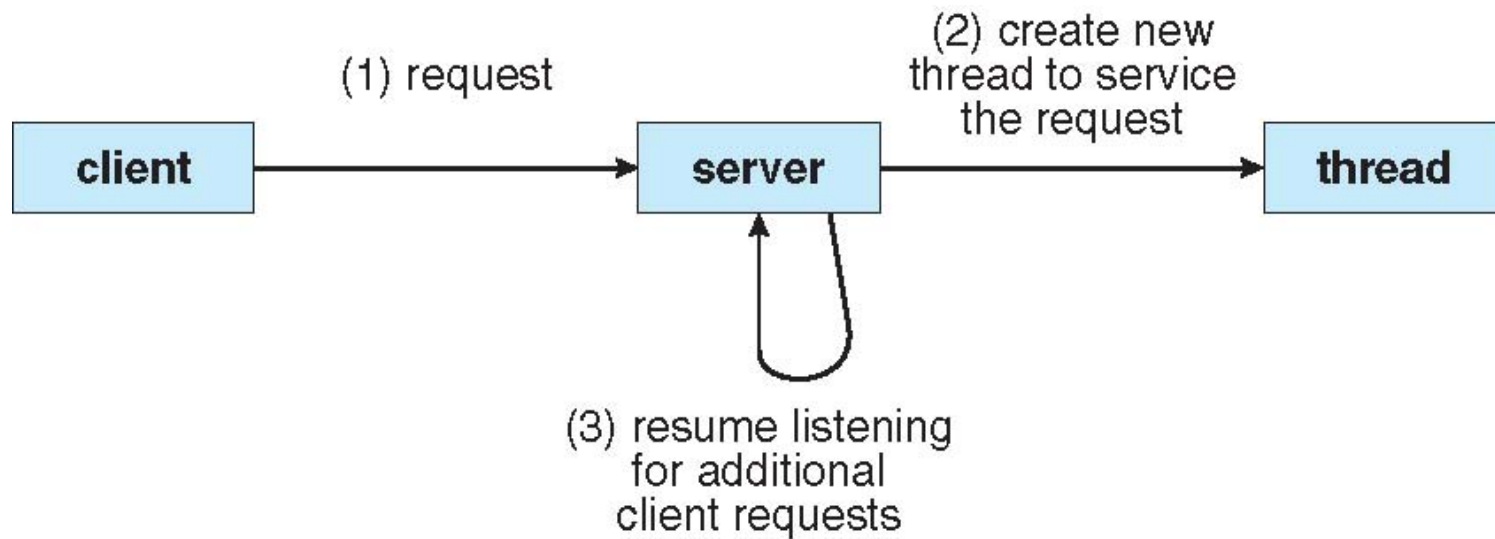
Threads vs. Processes

- Why multiple threads? Why not multiple processes?
 - Yes, all roads lead to Rome
 - But no, all roads are not equal

Threads vs. Processes

- Advantages of multi-threading over multi-processes
 - Far less time to create/terminate thread than process
 - An order of magnitude improvement
 - Less states to create/copy/inherit, or to release
 - Context switch is quicker between threads of the same process
 - Communication btw. threads of the same process is more efficient
 - Through shared memory
 - Makes complex program control structure simpler
 - Blocking I/O vs. non-blocking I/O

Multithreaded Server Architecture



Threads vs. Processes: An Example

```
main()
{
    sock = socket( ... );
    listen( ... );
    while (1) {
        new_sock = accept (sock);

        if ( fork()== 0 ){
            http_service(new_sock);
            exit(0);
        }
    }
}

http_service(int s){
    /*serve HTTP request & exit*/
    ...
}
```

```
main()
{
    sock = socket( ... );
    listen( ... );
    while (1) {
        new_sock = accept (sock);

        pthread_create(&tid, NULL,
http_service, new_sock);
    }
}

http_service(int s){
    /*serve HTTP request & exit*/
    ...
}
```

When to, and not to use threads?

- Applications
 - Multiprocessor machines
 - Handle slow devices
 - Background operations
 - Windowing systems
 - Server applications to handle multiple requests
- No threads cases
 - When each unit of execution require different authentication/user-id
 - E.g., secure shell server



Thread context switch

- A context switch takes place when
 - Async. events (e.g., time quota expires, hardware interrupt)
 - Sync. call
 - Thread performs a blocking system call
 - Thread voluntarily yields, e.g., lock unavailable
- Save state of currently executing thread
 - Copy all user registers to thread control block
 - Save control information: PC, SP
- Restore state of thread to run next
 - Copy values of user/control registers from thread control block to processor registers

Thread Implementation: User vs. Kernel level

- User-level threads
 - Thread management by user space thread library
 - Kernel not aware of thread activities
 - Thread library performs like an operating system
 - Thread creation & termination
 - Thread scheduling and context switches
 - Maintains control and context information
 - State, priority
 - User registers, stack, stack pointers
 - Can be preemptive or non-preemptive
- Kernel-level threads (lightweight processes)
 - Kernel supports multiple execution contexts (kernel threads)
 - Thread management done by the kernel

User level vs. Kernel level Threads

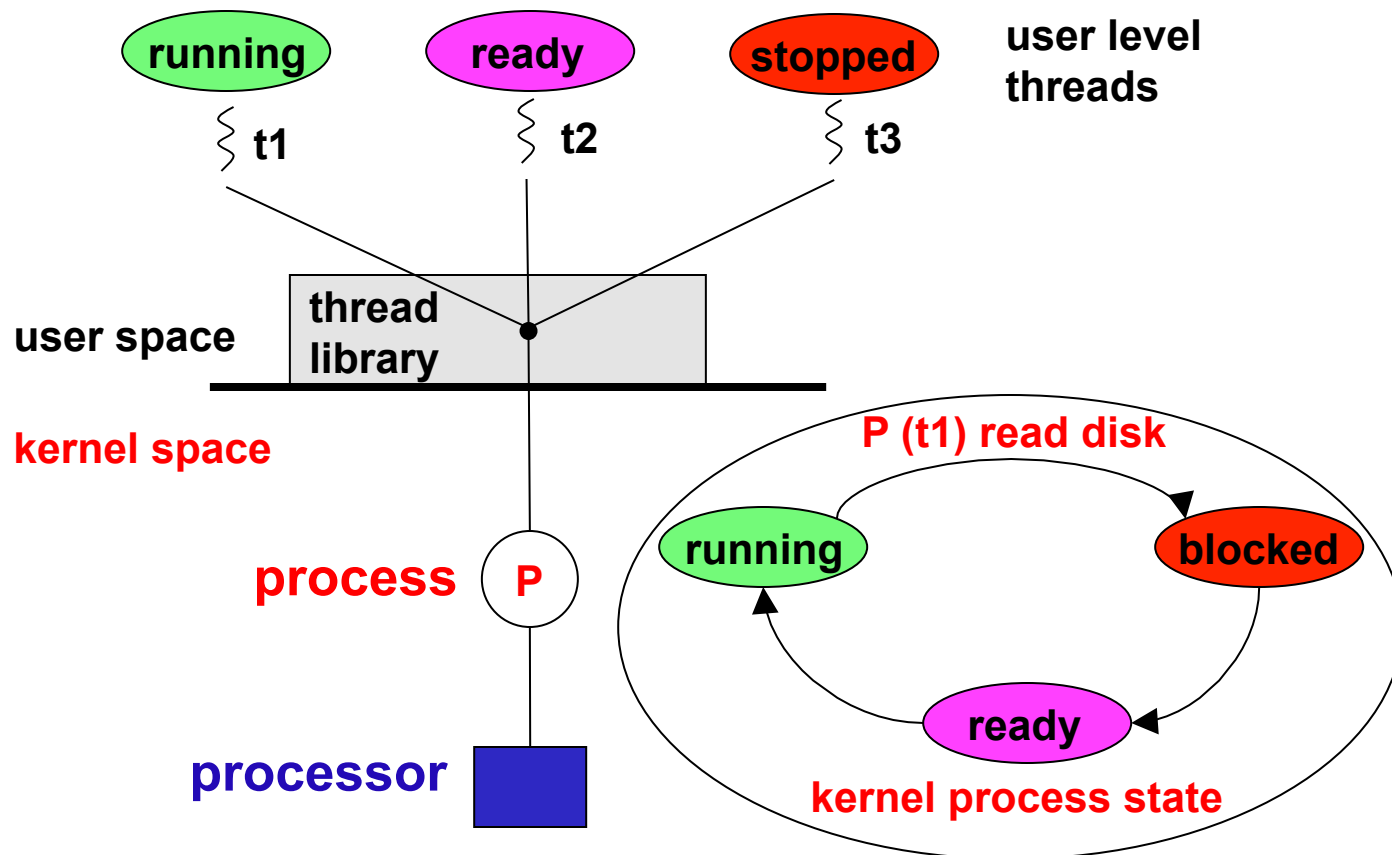


User Level Threads

- Advantages of user-level threads
 - Performance: low-cost thread operations and context switching
 - No kernel involvement, saves user/kernel mode switches
 - VAX results: $7\ \mu\text{s}$ (procedure call) vs. $17\ \mu\text{s}$ (kernel trap)
 - Flexibility: scheduling can be application-specific w/o kernel change
 - Portability: user-level thread library easy to port
- Disadvantages of user-level threads
 - If a user-level thread is blocked in the kernel, the entire process (all threads of that process) are blocked
 - Cannot take advantage of multiprocessing (the kernel assigns one process to only one processor)

User Level Thread: Blocking Problem

- Thread t1 issues blocking disk read
- Process **p** suspended by kernel
- Although thread t2 is ready to run, kernel not aware of user threads, and all threads are blocked

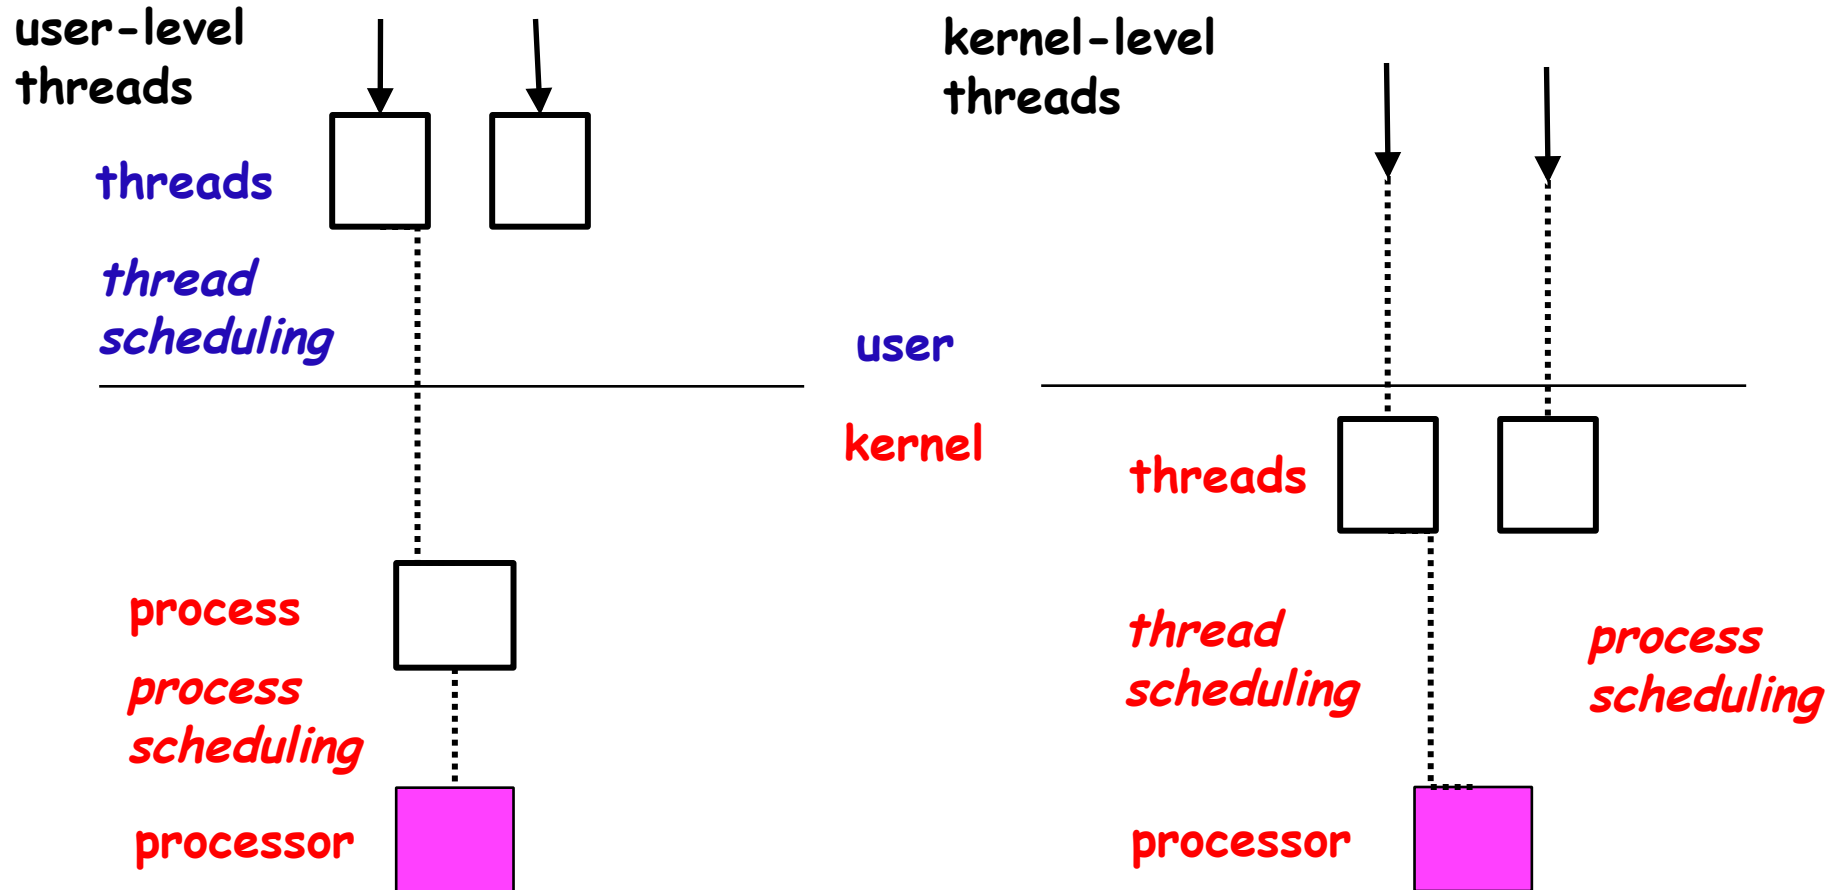


Kernel Level Threads

- Advantages of kernel-level threads
 - Blocking of one thread does not block other threads in a process
 - Threads within one process can be assigned to different processors simultaneously, taking advantage of SMP
- Disadvantages of kernel-level threads
 - Context switch btw. threads of one process requires two user/kernel mode switches
 - Inflexible in scheduling: entirely relies on a generic kernel scheduler
 - Portability: OS has to provide support

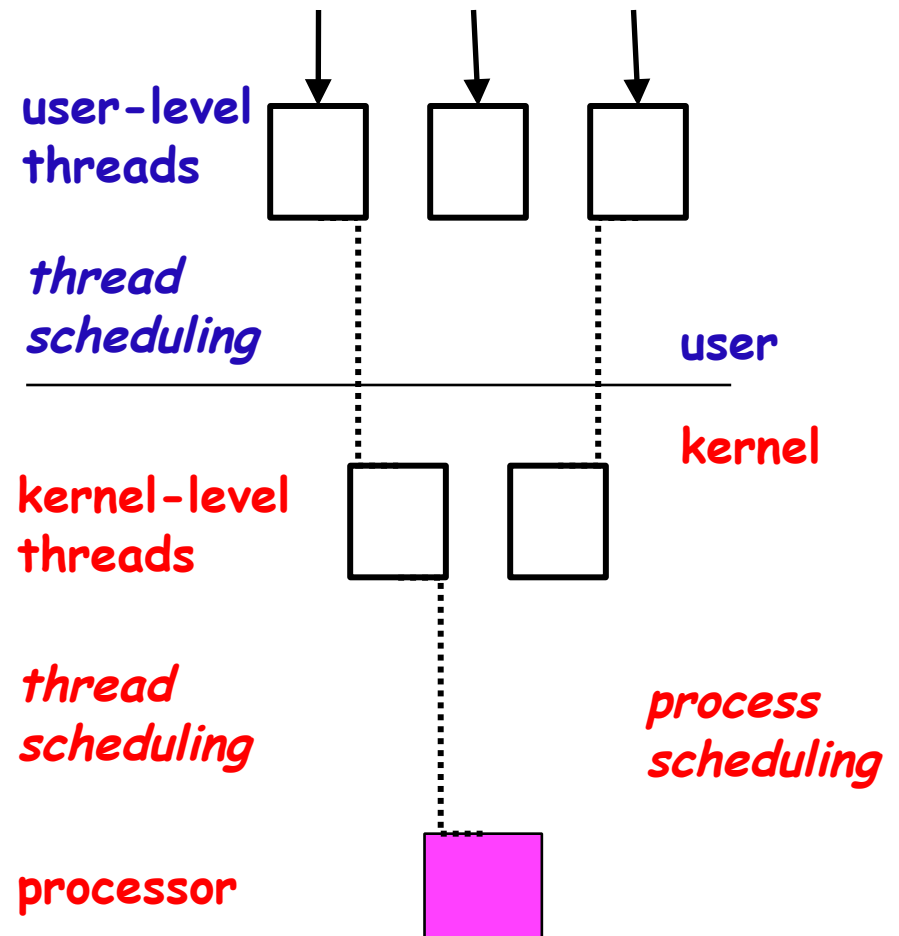
Operation	User thread (μ s)	Kernel thread (μ s)	Processes (μ s)
Null fork	34	948	11,300
Signal-wait	37	441	1,840

User-Level vs. Kernel-Level Threads



User-Level vs. Kernel-Level Threads

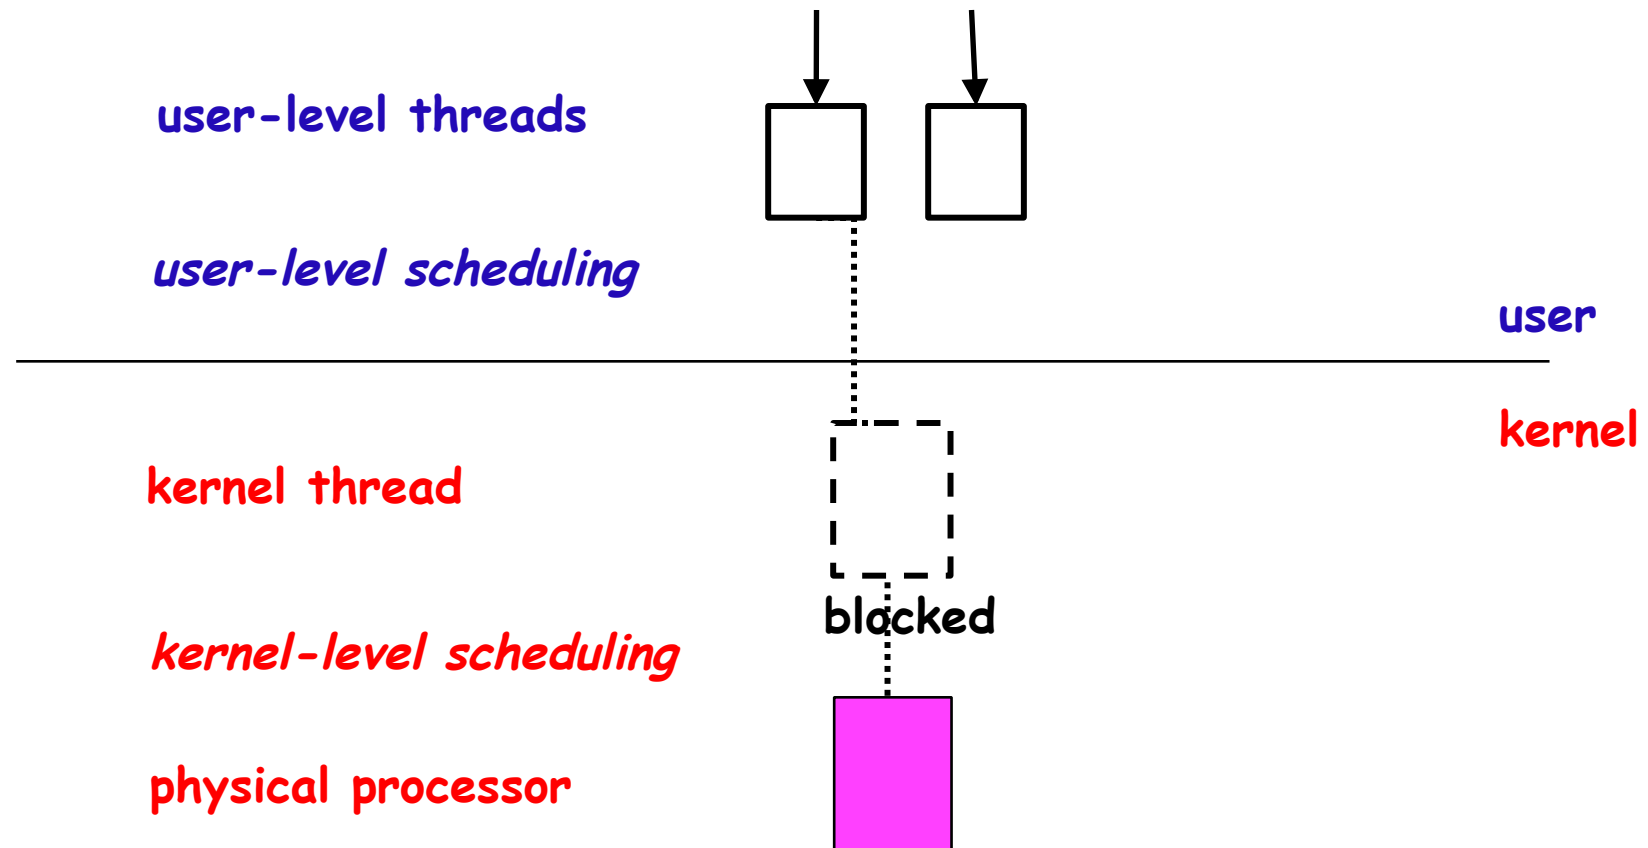
- No reason why we shouldn't have both (e.g. Solaris)
- Most systems now support kernel threads
- User-level threads are available as linkable libraries



Kernel Support for User-Level Threads

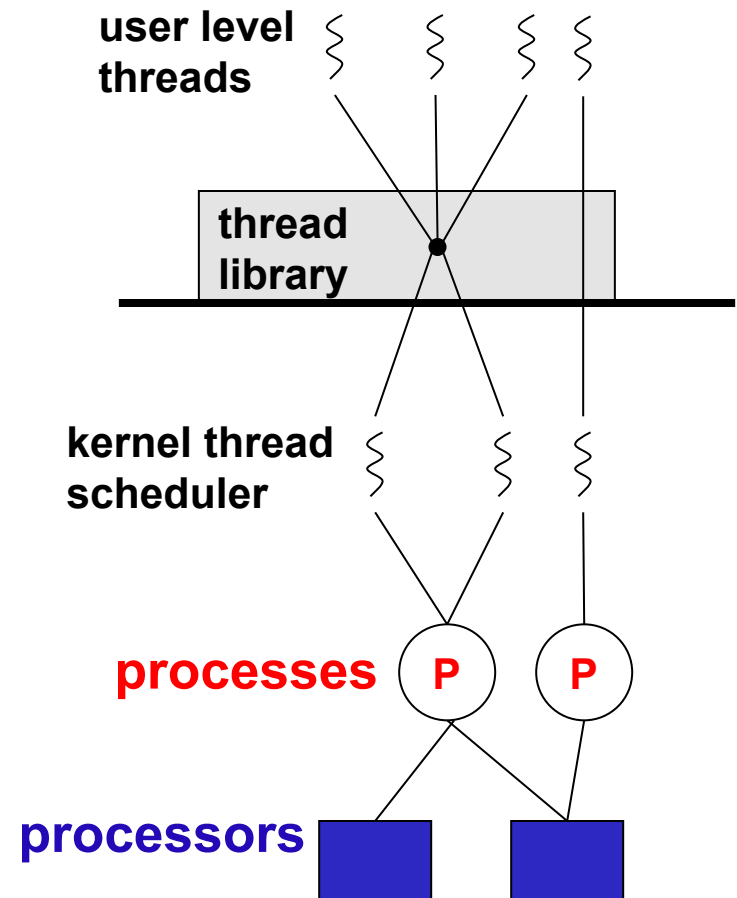
- Even kernel threads are not quite the right abstraction for supporting user-level threads
- Mismatch between where the scheduling information is available (user) and where scheduling on real processors is performed (kernel)
 - When the kernel thread is blocked, the corresponding physical processor is lost to all user-level threads although there may be some ready to run.
 - Scheduler Activation

Why Kernel Threads Are Not The Right Abstraction



Combining the user & kernel level threads

- Hybrid approach (e.g. Solaris)
- User level thread library
 - Thread creation/termination
 - Most of scheduling and synchronization
- User threads mapped to a set of kernel threads
 - Programmer can adjust how the mapping is done
 - 1 user thread – 1 kernel thread
 - m user thread – 1 kernel thread
 - m user thread – n kernel thread




Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many
- Two-Level


Many-to-One

- Many user-level threads mapped to single kernel thread
- Examples
 - Solaris Green Threads
 - GNU Portable Threads




One-to-One

- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later




Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the *ThreadFiber* package

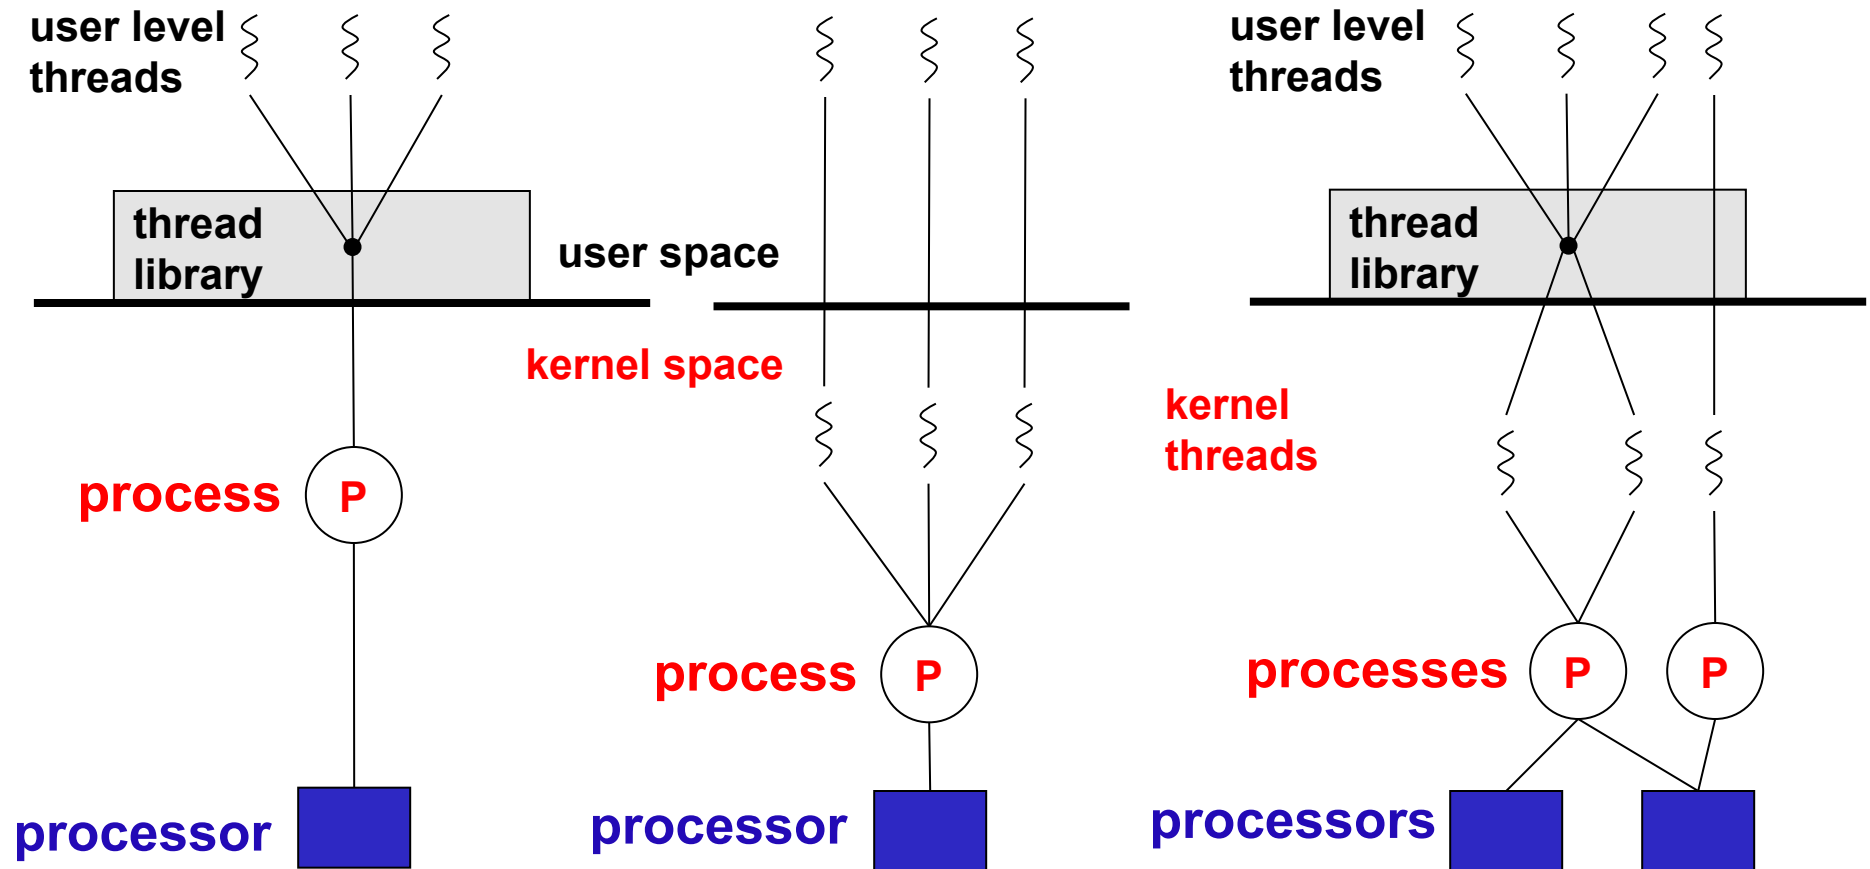


Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



User level, Kernel level, and Hybrid Threads



Signal Handling in Multithreaded Process

- Signal is process-level concept
 - Handler associated with the process
 - When a thread register a signal handler, it applies to all threads in the same process
 - Signal is only delivered to a single thread, to avoid multiple delivery
 - A thread, however, can ignore/mask certain signals
- Signal delivery
 - What if all threads are blocked when event occurs?
 - OS writes into PCB: that a signal should be delivered to the process
 - Next time when a thread that does not mask the signal is dispatched, the signal is delivered
- Common strategy:
 - One thread responsible to field all signals
 - calls `sigwait()` in an infinite loop
 - All other threads mask all signals

Thread Cancellation

- Terminating a thread before it has finished
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled