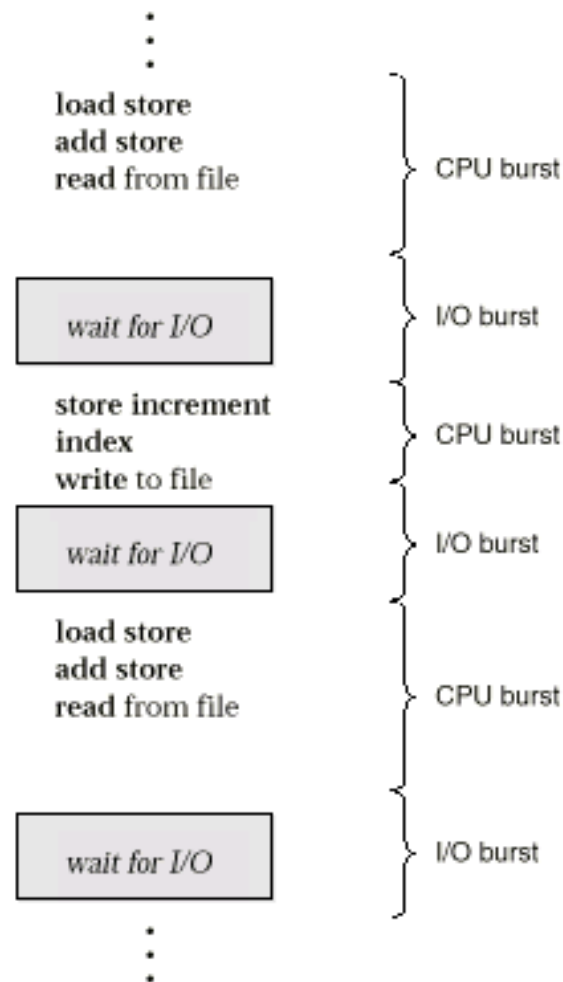



CPU Scheduling

Process Execution: Alternating Sequence of CPU And I/O Bursts



Histogram of CPU-burst Times



CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Terminates
 3. Switches from running to ready state
 4. Switches from waiting to ready
- Scheduling under 1,2 is *nonpreemptive*
- Scheduling under 3,4 is *preemptive*

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- *Dispatch latency* – time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

- CPU utilization – keep the CPU as busy as possible
- Throughput – # of processes that complete their execution per time unit
- Turnaround time – amount of time to execute a particular process
- Waiting time – amount of time a process has been waiting in the ready queue
- Response time – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

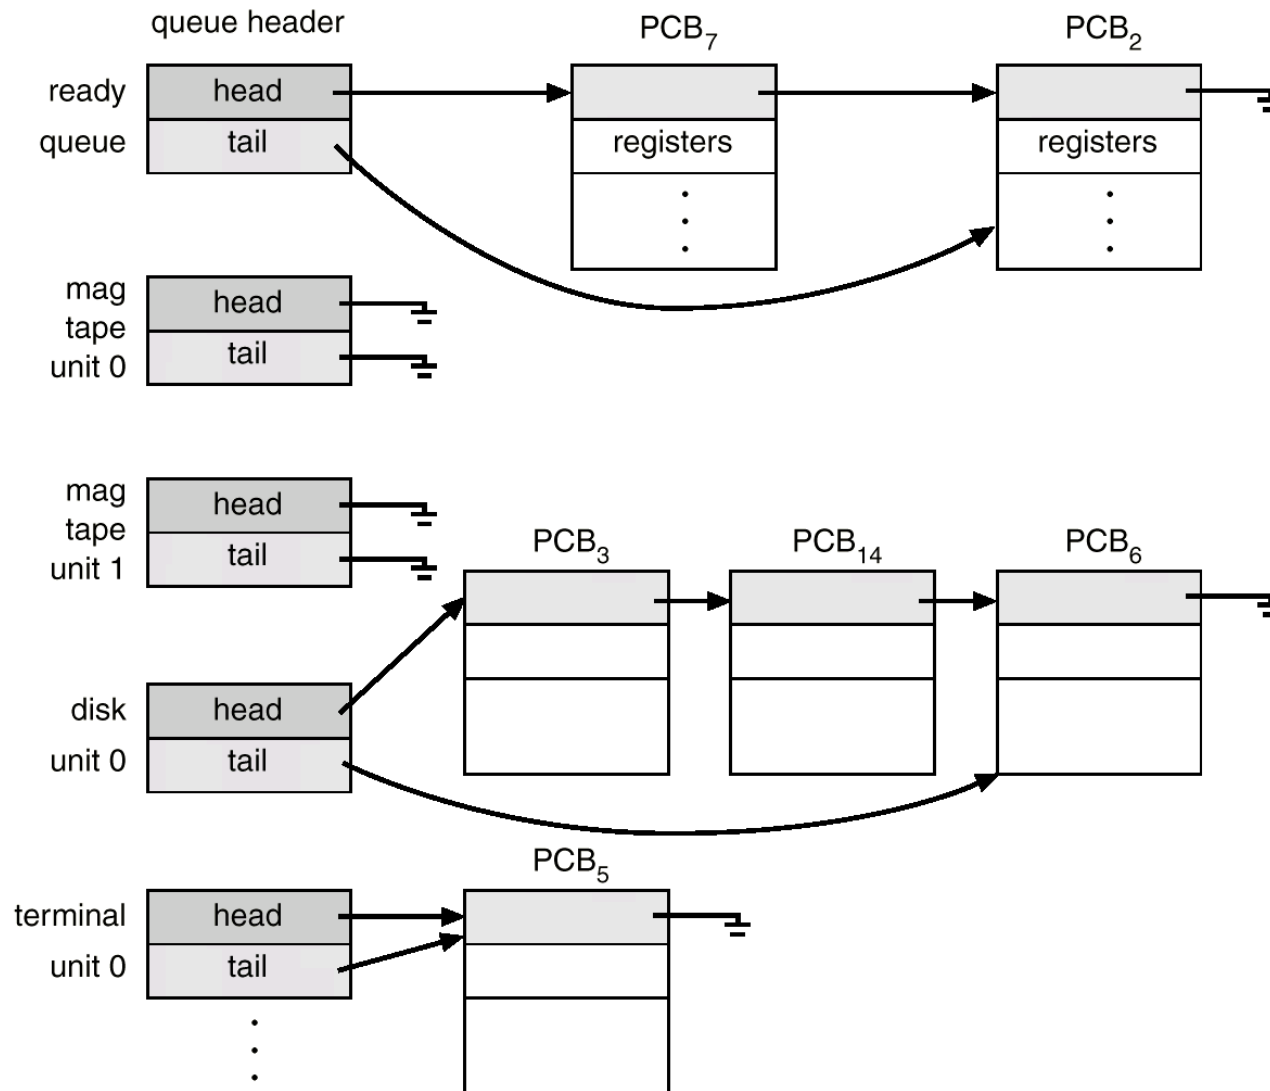
Conflicting criteria

- Minimizing response time requires more context switches for many processes
 - → incur more scheduling overhead
 - → decrease system throughput
- Scheduling algorithm depends on nature of system
 - Batch vs. interactive
 - Designing a generic AND efficient scheduler is difficult


Process Scheduling Queues

- **Ready** queue
 - Set of processes residing in main memory, ready, and waiting to execute
- **Job** queue
 - Set of all processes in the system
- **Device** queues
 - Set of processes waiting for an I/O device
- Process migration between the various queues

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling



Schedulers

- **Long-term** scheduler (or job scheduler)
 - Which processes should be brought into the ready queue
 - Invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
 - Controls the **degree of multiprogramming**
- **Short-term** scheduler (or CPU scheduler)
 - Which process should execute next (allocates CPU)
 - Invoked very frequently (milliseconds) \Rightarrow (must be fast)

Schedulers (Cont.)

- Processes can be described as either:
 - *I/O-bound process*
 - spends more time doing I/O than computations
 - many short CPU bursts
 - *CPU-bound process*
 - spends more time doing computations
 - few very long CPU bursts

Scheduling Algorithms

- Make scheduling decisions
 - Which process(es) to run
 - For how long
 - When to swap running process out
- Examples
 - First come first serve (FCFS), round robin, shortest job first (SJF)
 - Modern OS: priority-based algorithms

Example Scheduling in XINU

- Each process assigned a *priority*
 - Non-negative integer value
 - Initialized when process created
 - Can be changed
- Scheduler chooses process with the highest priority
 - Processes with the same priority are scheduled in a round-robin fashion
- Policy enforced as a system-wide invariant

Implementation of Scheduling

- Process eligible if state is
 - *ready or current*
- To avoid searching process table
 - Keep ready processes on linked list called *ready list*
 - Order ready list by priority
 - Selection in constant time

Example Scheduler Code

```
int resched()
{
    register struct pentry *optr; /* pointer to old process entry */
    register struct pentry *nptr; /* pointer to new process entry */

    /* no switch needed if current process priority higher than next */
    if ( ( (optr= &proctab[currpid])->pstate == PRCURR) &&
        (lastkey(rdytail)<optr->pprio) ) {
        return(OK);
    }

    /* force context switch */
    if (optr->pstate == PRCURR) {
        optr->pstate = PRREADY;
        insert(currpid,rdyhead,optr->pprio);
    }

    /* remove highest priority process at end of ready list */
    nptr = &proctab[ (currpid = getlast(rdytail)) ];
    nptr->pstate = PRCURR; /* mark it currently running */
#ifdef RTCLOCK
    preempt = QUANTUM; /* reset preemption counter */
#endif
    ctxsw((int)&optr->pesp, (int)optr->pirmask, (int)&nptr->pesp, (int)nptr->
    >pirmask);

    /* The OLD process returns here when resumed. */
    return OK;
}
```

Puzzle #1

- Invariant says that at any time, one process must be executing
- Context switch code moves from one process to another
- Question: which process executes the context switch code?

Solution to Puzzle #1

- “Old” process
 - Executes first half of context switch
 - Is suspended
- “New” process
 - Continues executing where previously suspended
 - Usually runs second half of context switch

Puzzle #2

- Invariant says that at any time, one process must be executing
- All user processes may be idle (e.g., applications all wait for input)
- Which process executes?

Solution to Puzzle #2

- OS needs an extra process
 - Called *NULL process*
 - Never terminates
 - Cannot make a system call that takes it out of ready or current state
 - Typically an infinite loop

Lab 1 – Process Scheduling

- *Revisit Xinu scheduling invariant:*

At any time, the CPU must run the highest priority eligible process. Among processes with equal priority, scheduling is round robin

- *Question: what is a potential problem here?*

PA 1 – Process Scheduling

- The process scheduling policy has a limitation, namely process **starvation**
- You are asked to implement two different policies
 - **Random scheduler**
 - **Linux Scheduling**

Random Scheduler

- Total N processes P_i ($i=0..N-1$) in the ready queue
- Each P_i
 - Priority: $PRIO_i$
 - Generate a random number x between $[0, \text{sum}(PRIO_i)-1]$
 - If $x < \text{the highest priority}$, pick the first process in the queue
 - Otherwise, $x = x - \text{the highest priority}$, and check whether $x < \text{the second highest priority}$ and so on.

Linux Like Scheduling

- Epoch based scheduling
- Dynamically adjust per-process quantum at the beginning of epoch
 - Quantum defines how many CPU ticks are allocated to the process
 - Consider both priority and past usage
 - E.g. $\text{quantum} = \text{priority} + \text{unused CPU ticks in previous epoch} / 2$
- Select the process with the highest goodness
 - Goodness = 0 if the process uses up its quantum
 - Goodness = priority + unused CPU ticks
 - Round robin among equal goodness

Unix Scheduling – 4.3 BSD

- Multilevel feedback
- For process j at the beginning of timer interval i
 - $P_j(i) = \text{Base}_j + \text{CPU}_j(i) + \text{Nice}_j$
- $\text{CPU}(i)$: an estimate of *recent* CPU usage
 - More CPU use leads to lower priority
 - Uses current load (number of ready processes)
 - Aging: $\text{CPU}(i) = \text{CPU}(i-1)/2 + \text{load}(i)/2$
- $P_j(i)$: priority of j at beginning of interval i ; lower value \rightarrow higher priority
- Base_j : base priority of Process j
 - Divided into bands: Swapper, Block I/O device, File, I/O, User process
- Nice_j : user controllable factor

Linux Scheduling

- Two algorithms: time-sharing and real-time
- Time-sharing
 - Prioritized credit-based – process with most credits is scheduled next
 - Credit subtracted when timer interrupt occurs
 - When credit = 0, another process chosen
 - When all processes have credit = 0, recrediting occurs
 - Based on factors including priority and past CPU usage
- Real-time
 - Soft real-time
 - Rate monotonic scheduling (not covered)
 - End deadline first scheduling (not covered)

PA 1 – Process Scheduling

- Read relevant source code in Xinu
 - Process queue management
 - `h/q.h` `sys/queue.c` `sys/insert.c`, ...
 - Proc. creation/suspension/resumption/termination:
 - `sys/create.c`, `sys/suspend.c` `sys/resume.c`, `sys/kill.c`
 - Priority change
 - `sys/chprio.c`
 - Process scheduling
 - `sys/resched.c`
 - Other initialization code
 - `sys/initialize.c`

Dynamic Priority-based Scheduling

- In Xinu: Timer interrupt handler
 - Related files: `sys/clkint.S` `sys/clkinit.c`
 - Interrupt rate – based on clock timer
 - `ctr1000`: 1ms
 - Scheduling rate:
 - $\text{Interrupt rate} * \text{QUANTUM}$
 - Others
 - `preempt`: preemption counter

Next Lecture

- Process Synchronization