

Deadlock

The Deadlock Problem

- Definition
 - A set of blocked processes each holding a resource and waiting to acquire a resource held by another process
 - None of the processes can proceed or back-off (release resources it owns)
- Example
 - semaphores A and B , initialized to 1

P_0	P_1
wait (A);	wait(B)
wait (B);	wait(A)

Deadlock Conditions

- Deadlock can arise if four conditions hold simultaneously
 - **Mutual exclusion:** only one process at a time can use a resource instance
 - **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
 - **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
 - **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Methods for Handling Deadlocks

- **Ignore** the problem and pretend that deadlocks would never occur.
- **Prevent** the system from entering a deadlock state.
- Allow the system to enter a deadlock state and then **detect/recover**.

The IGNORE Approach

- Pretend there is no problem
 - Unfortunately they can occur
 - Reasonable if
 - Deadlocks occur very rarely and cost of prevention is high
- Do your typical OSES take this approach?
- It is a trade off between
 - Overhead
 - Correctness

The PREVENT Approach

- Restrain the ways requests can be made to break one of the four necessary conditions for deadlocks
- Attacking the **mutual exclusion** condition:
 - Some devices (such as printer) can be spooled
 - only the printer daemon uses printer resource
 - thus deadlock for printer eliminated
 - Not all devices can be spooled

The PREVENT Approach

- Attacking the **Hold and Wait** Condition:
 - Require processes to request all resources before starting
- Problems
 - may not know required resources at start of run
 - also ties up resources other processes could be using
- Variation:
 - before a process requests for a new resource, it must give up all resources and then request all resources needed

The PREVENT Approach

- Attacking the **No Preemption** Condition:
 - When a process holding some resources and waiting for others, its resources may be preempted to be used by others
- Problems
 - Many resources may not allow preemption; i.e., preemption will cause process to fail

The PREVENT Approach


- Attacking the Circular Wait Condition:
 - Impose a total order of all resource types; and require that all processes request resources in the same order

Deadlock Avoidance

- When a process requests available resource, system must decide if immediate allocation leaves the system in a **safe** state
- System is in **safe** state if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of all processes, such that
 - For each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
 - That is:
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Deadlock Avoidance

- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.






Resource-Allocation Graph


- A set of vertices V and a set of edges E .
 - V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the **processes** in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all **resource types** in the system
 - E is partitioned into two types:
 - **request edge** – directed edge $P_i \rightarrow R_j$
 - **assignment edge** – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph

- Process
- Resource type with 4 instances
- P_i requests instance of R_j
- P_i is holding an instance of R_j




Resource Allocation Example 1




- Is there deadlock?

Resource Allocation Graph Example 2



- Is there a deadlock?

Resource Allocation Graph Example 3



- Is there a deadlock?

Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.

Resource-Allocation Graph Algorithm

- Suppose that process P_i requests a resource R_j
- The request can be granted **only if**
 - Converting the request edge to an assignment edge does not result in the formation of a cycle in the resource-allocation graph

Banker's Algorithm

- Each process must a priori claim the **maximum** set of resources that might be needed in its execution.
- Safety check
 - Repeat
 - pick any process that can finish with existing available resources; finish it and release all its resources
 - until no such process exists
 - all finished → safe; otherwise → unsafe.

Data Structure for the Banker's Algorithm

- $n = \#$ of processes, $m = \#$ of resources types.
 - **Available:** Vector of length m
 - If available $[j] = k$, there are k instances of resource type R_j available
 - **Max:** $n \times m$ matrix.
 - If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j
 - **Allocation:** $n \times m$ matrix.
 - If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j
 - **Need:** $n \times m$ matrix.
 - If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$

Safety Algorithm

- Step 1: Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
 - *Work* = *Available*
 - *Finish* [i] = *false* for $i = 0, 1, \dots, n-1$
- Step 2: Find any i such that both (If no, Step 4)
 - *Finish* [i] = *false*
 - $Need_i \leq Work$
- Step 3. $Work = Work + Allocation_i$
 - *Finish* [i] = *true*
 - Step 2
- Step 4. If *Finish* [i] == *true* for all i , then the system is in a safe state

Resource-Request Algorithm for Process P_i

- Process P_i wants k instances of R_j ($Request_i[j] = k$)
 - Step 1: If $Request_i \leq Need_i$, go to step 2
 - Otherwise, raise error condition, since process has exceeded its maximum claim
 - Step 2: If $Request_i \leq Available$, go to step 3
 - Otherwise P_i must wait, since resources are not available
 - Step 3: Pretend to allocate requested resources to P_i by modifying the state as follows:

$Available = Available - Request_i;$
 $Allocation_i = Allocation_i + Request_i;$
 $Need_i = Need_i - Request_i;$

 - If safe \Rightarrow the resources are allocated to P_i
 - If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example of Banker's Algorithm

- 5 processes P_0 through P_4 ;
- 3 resource types:
 - A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Question:

Is this a safe state?

Question:

Can request for (1,0,2) by P_1 be granted?

Example of Banker's Algorithm

- 5 processes P_0 through P_4 ;
- 3 resource types:
 - A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

Question:

Can request for (3,3,0) by P_4 be granted?

Question:

Can request for (0,2,0) by P_0 be granted?


Methods for Handling Deadlocks

- **Ignore** the problem and pretend that deadlocks would never occur.
- **Prevent** the system from entering a deadlock state.
- Allow the system to enter a deadlock state and then **detect/recover**.

Single Instance of Each Resource Type


- Maintain *wait-for* graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph

Single Instance of Each Resource Type



(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph

Additional Issues

- When there are several instances of a resource type
 - cycle detection in wait-for graph is not sufficient.
- Deadlock detection is very similar to the safety check in the Banker's algorithm

Recovery from Deadlock

- Recovery through preemption
 - take a resource from some other process
 - depends on nature of the resource
- Recovery through rollback
 - checkpoint a process state periodically
 - rollback a process to its checkpoint state if it is found deadlocked
- Recovery through killing processes
 - kill one+ of the processes in the deadlock cycle
 - the other processes get its resources
 - In which order should we choose process to kill?