

# Interrupts

# Introduction

- Interrupts provide an efficient way to handle **unanticipated events** and improve **processor utilization**
- Interrupts alter a program's flow of control
  - Interrupt causes transfer of control to an *interrupt service routine* (ISR)
    - ISR is also called a *handler*
  - When the ISR is completed, the original program resumes execution
  - Behavior is similar to a procedure call
    - Some significant differences between the two

# Interrupts vs. Procedures

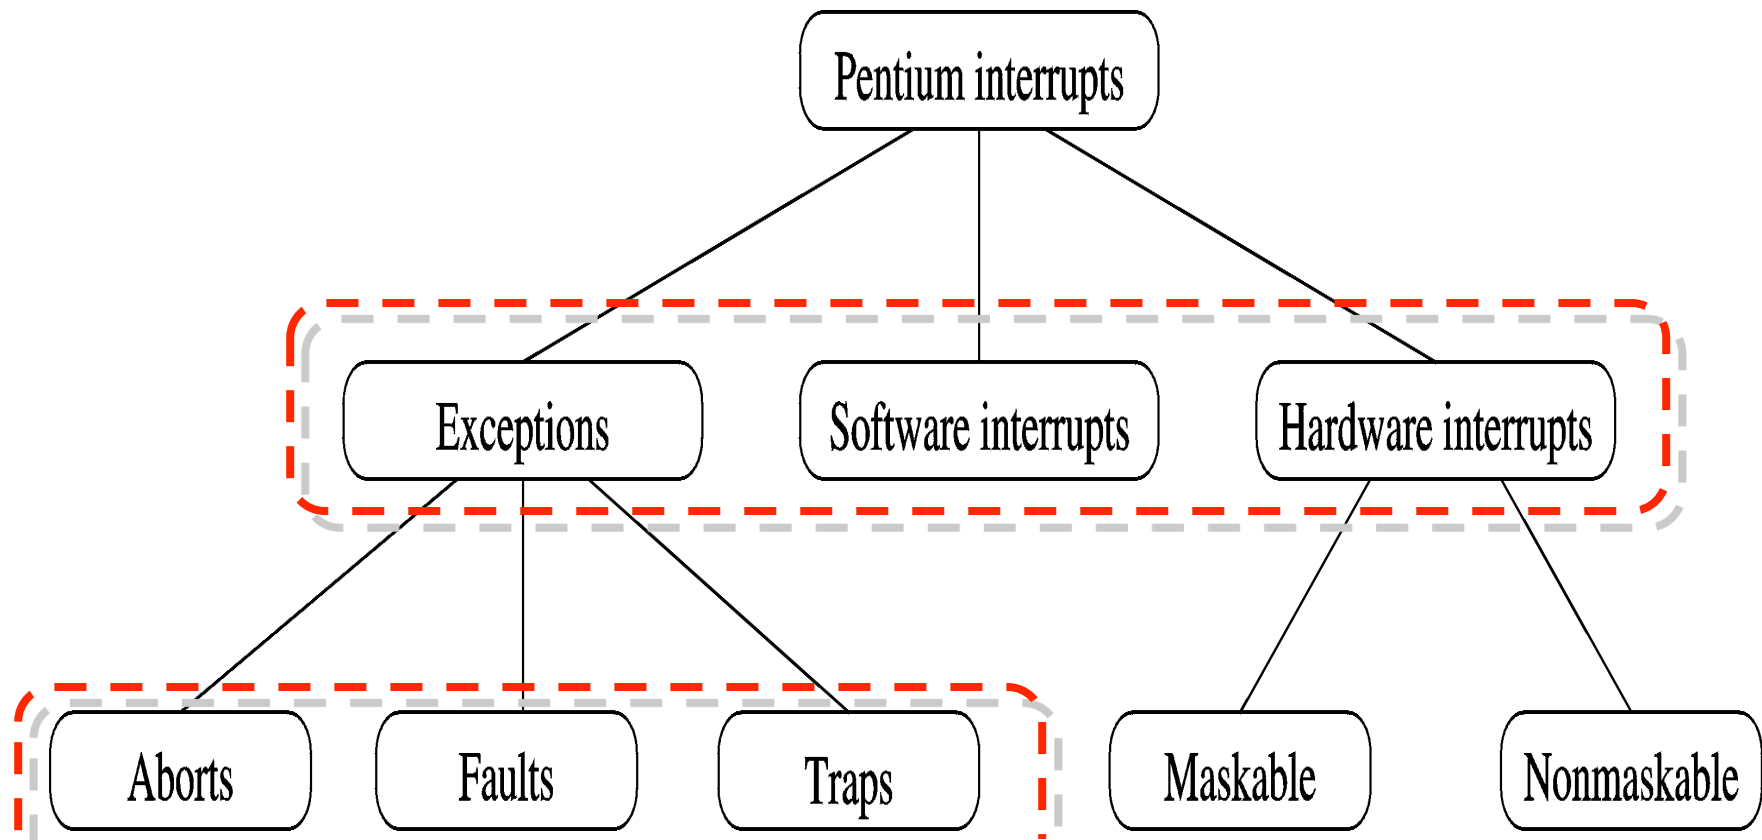
## Interrupts

- Initiated by both *software* and *hardware*
- Can handle *anticipated* and *unanticipated* internal as well as external events
- ISRs or interrupt handlers are memory resident
- Use numbers to identify an interrupt service
- **eflags** register is saved automatically

## Procedures

- Can only be initiated by *software*
- Can handle *anticipated* events that are coded into the program
- Typically loaded along with the program
- Use meaningful names to indicate their function
- Do not save the **eflags** register

# A Taxonomy of Pentium Interrupts



## Difference:

- ☐ Depending on the way they are reported
- ☐ Whether or not the interrupted instruction is restarted

# Exceptions: Faults, Traps, and Aborts

- Faults
  - Instruction boundary before the instruction during which the exception was detected
  - Restarts the instruction
- Examples:
  - Page fault
  - Segment-not-found fault

# Exceptions: Faults, Traps, and Aborts

- Traps
  - Instruction boundary immediately after the instruction during which the exception was detected
  - No instruction restart
- Examples:
  - Overflow exception (interrupt 4) is a trap
  - User defined interrupts are also examples of traps

# Exceptions: Faults, Traps, and Aborts

- Aborts
  - No precise location of the instruction that caused the exception
  - No instruction restarting
  - Reporting severe errors such as hardware errors and inconsistent values in system tables
- Examples:
  - Machine check
  - Double fault

# Dedicated Interrupts

- Several Pentium predefined interrupts --- called dedicated interrupts
- These include the first five interrupts:

| <b>interrupt type</b> | <b>Purpose</b>               |
|-----------------------|------------------------------|
| 0                     | Divide error                 |
| 1                     | Single-step                  |
| 2                     | Non-maskable interrupt (NMI) |
| 3                     | Breakpoint                   |
| 4                     | Overflow                     |



## Dedicated Interrupts (cont'd)

- Single-Step Interrupt
  - Useful in debugging
  - To single step, Trap Flag (TF) should be set
  - CPU automatically generates a type 1 interrupt after executing each instruction if TF is set
  - Type 1 ISR can be used to present the system state to the user

## Dedicated Interrupts (cont'd)

- Breakpoint Interrupt
  - Useful in debugging
  - CPU generates a **type 3** interrupt
  - Generated by executing a special single-byte version of **int 3** instruction (opcode CCH)

# Interrupt Taxonomy

- Exceptions
- Software Interrupts
- Hardware Interrupts

# Software Interrupts

- Initiated by executing an *int* instruction, where the interrupt number is an integer between 0 and 255
- Each interrupt can be parameterized to provide several services.
  - For example, Linux interrupt service *int 0x80* provides a large number of services (more than 330 system calls!)
    - EAX register is used to identify the required service under *int 0x80*

# Hardware Interrupts

- Software interrupts are synchronous events
  - Caused by executing the `int` instruction
- Hardware interrupts are asynchronous in nature
  - Typically caused by applying an electrical signal to the processor chip
- Hardware interrupts can be
  - Maskable
  - Non-maskable

## How Are Hardware Interrupts Triggered?

- Maskable interrupt is triggered by applying an electrical signal to the **INTR** (INTerrupt Request) pin of Pentium
  - Processor recognizes this interrupt only if IF (interrupt enable flag) is set
  - Interrupts can be masked or disabled by clearing IF
- Non-maskable interrupt is triggered by applying an electrical signal to the **NMI** pin of processor
  - Processor always responds to this signal
  - Cannot be disabled under program control

## How Does the CPU Know the Interrupt Type?

- Interrupt invocation process is common to all interrupts
  - Whether originated in software or hardware
- For hardware interrupts, processor initiates an interrupt acknowledge sequence
  - processor sends out interrupt acknowledge (INTA) signal
  - In response, interrupting device places interrupt vector on the data bus
  - Processor uses this number to invoke the ISR that should service the device (as in software interrupts)

## How Can More Than One Device Interrupt?

- Processor has only one INTR pin to receive interrupt signal
- Typical system has more than one device that can interrupt --- keyboard, hard disk, etc.
- Use a special chip to prioritize the interrupts and forward only one interrupt to the CPU
  - 8259 Programmable Interrupt Controller chip performs this function



# Interrupt Processing

- How many interrupts can be supported?
  - Up to 256 interrupts
- Interrupt number is used as an index into the **Interrupt Descriptor Table** (IDT)
  - This table stores the addresses of all ISRs
  - Each descriptor entry is 8 bytes long
    - Interrupt number is multiplied by 8 to get byte offset into IDT

## Detailed Steps in Interrupt Processing

- Step 1: Save the current machine state
- Step 2: Load the machine state for interrupt handling
- Step 3: Invoke the corresponding ISR
- Step 4: Resume the program execution

### Question:

Why do we need to save the current machine states?

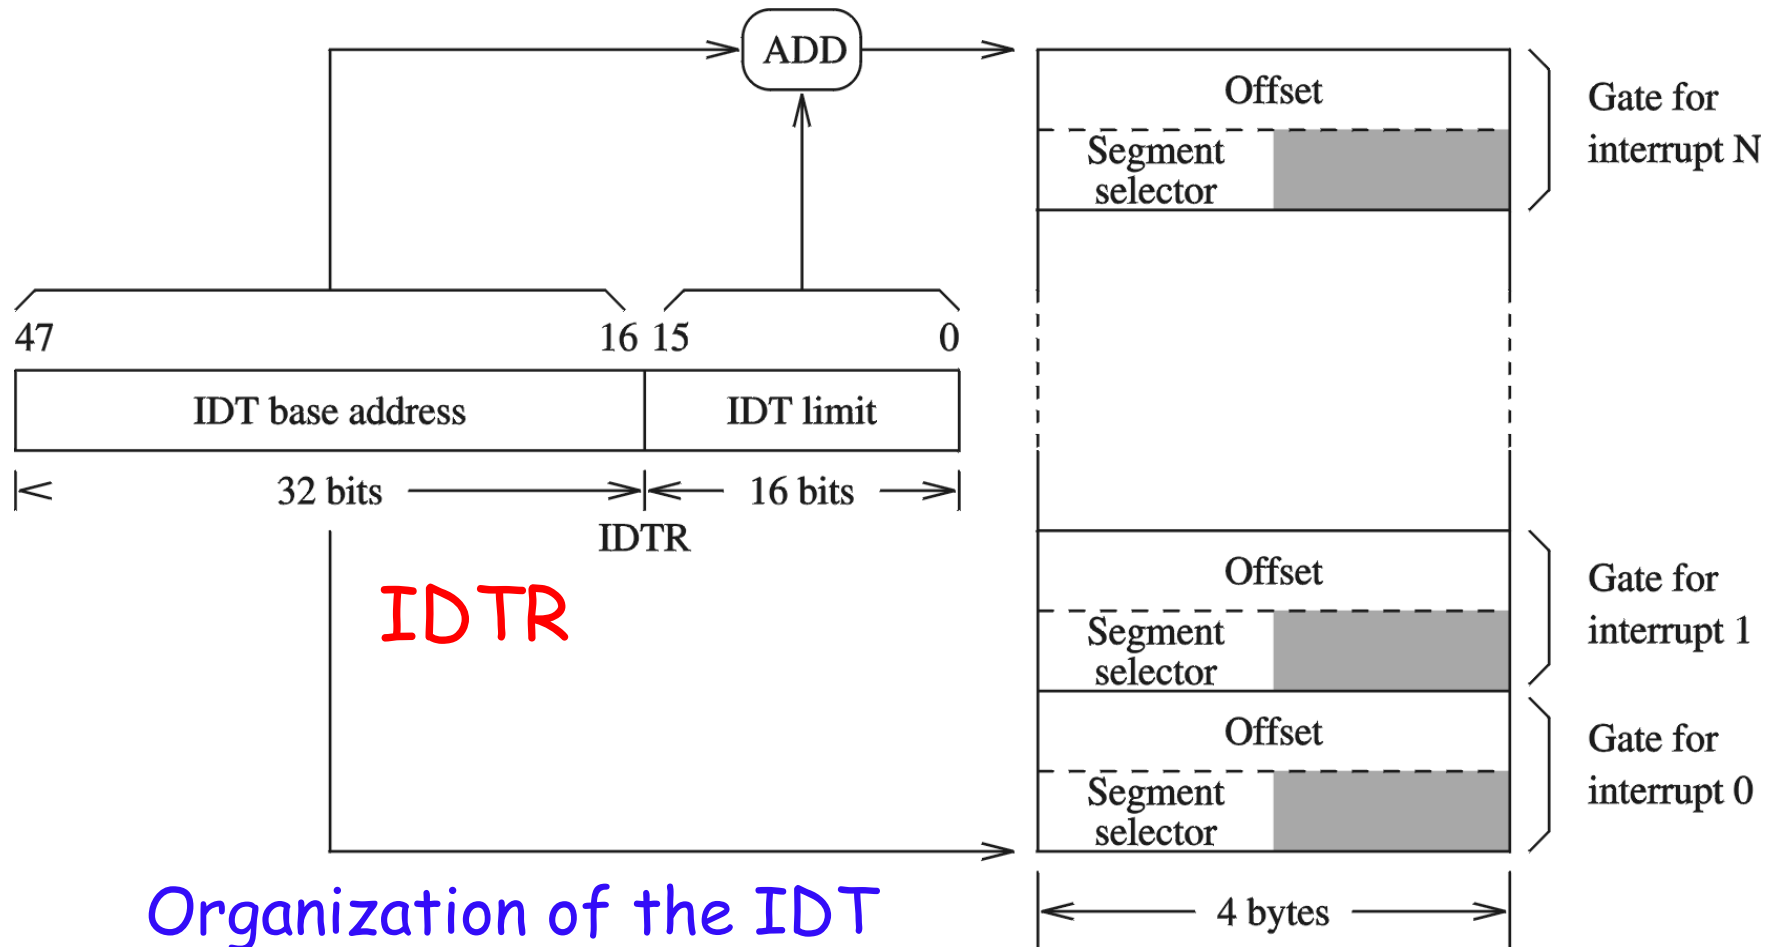
## Step 1: Save the Current Machine State

- Push the EFLAGS register onto the stack
- Clear interrupt enable and trap flags
  - This disables further interrupts
  - Use `cli` to clear interrupts
  - Use `sti` to enable interrupts
- Push CS and EIP registers onto the stack

## Step 2: Load the Machine State for Interrupt Handling

- Load CS (code segment register) with the 16-bit segment selector from the interrupt gate
- Load EIP (instruction pointer register) with the 32-bit offset value from the interrupt gate

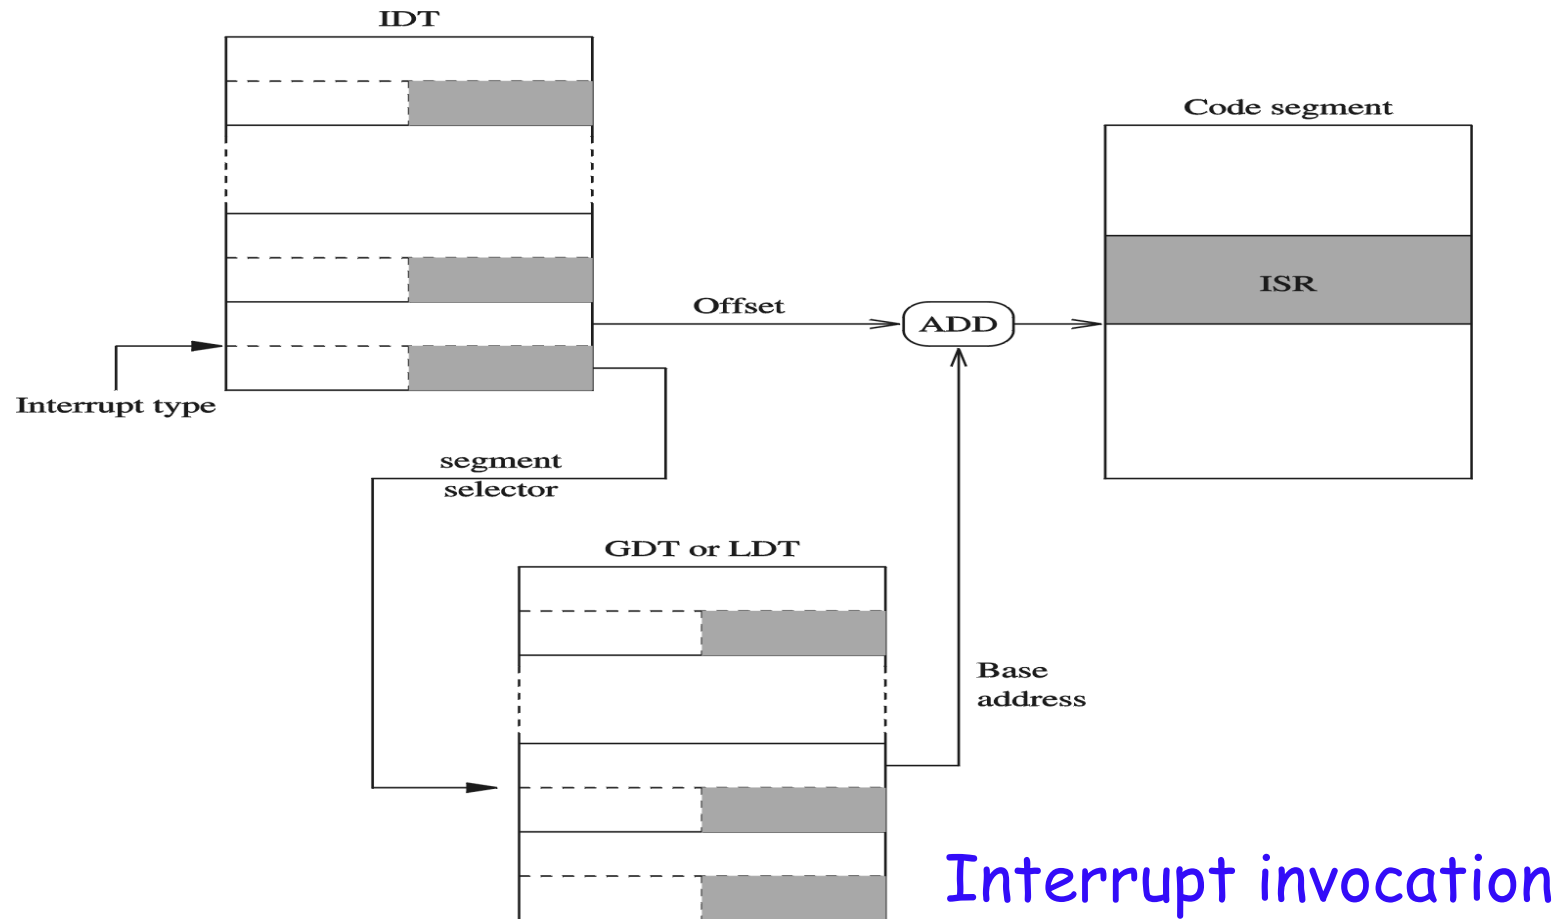
# Protected Mode Interrupt Processing



# Protected Mode Interrupt Processing

- IDTR contains the memory location of IDT
- IDTR is a 48-bit register
  - 32 bits for IDT base address
  - 16 bits for IDT limit value
    - IDT requires only 2048 (11 bits)
    - A system may have smaller number of descriptors
      - Set the IDT limit to indicate the size in bytes
- Two special instructions to load (`lidt`) and store (`sidt`) IDT
  - Both take the address of a 6-byte memory as the operand

# Protected Mode Interrupt Processing



## Step 3: Invoke the ISR

- ISR: Interrupt-specific service routine
- Examples:
  - Single-step
  - Breakpoint
  - Timer
  - Page fault
  - ...



## Step 4: Resume the Program Execution

- What is the last instruction in an ISR:
  - `iret`
- The actions taken on `iret` are:
  - pop the 32-bit value on top of the stack into EIP
  - pop the 16-bit value on top of the stack into CS
  - pop the 32-bit value on top of the stack into the EFLAGS register
- As in procedures, make sure that your ISR does not leave any data on the stack
  - Match your push and pop operations within the ISR

## An Example:

- Timer interrupt handler
  - Related files: `sys/clkint.S` `sys/clkinit.c`
  - Interrupt rate – based on clock timer
    - `ctr1000`: 1ms
  - Scheduling rate:
    - $\text{Interrupt rate} * \text{QUANTUM}$
- You will be familiar with page fault handler in PA 3!